

NNDL_ASSIGNMENT-7

NAME:JYOSHNA YARRAGUNTLA

STUDENT ID:700758848

Image Classification with CNN:

1. Training the model
2. Evaluating the model

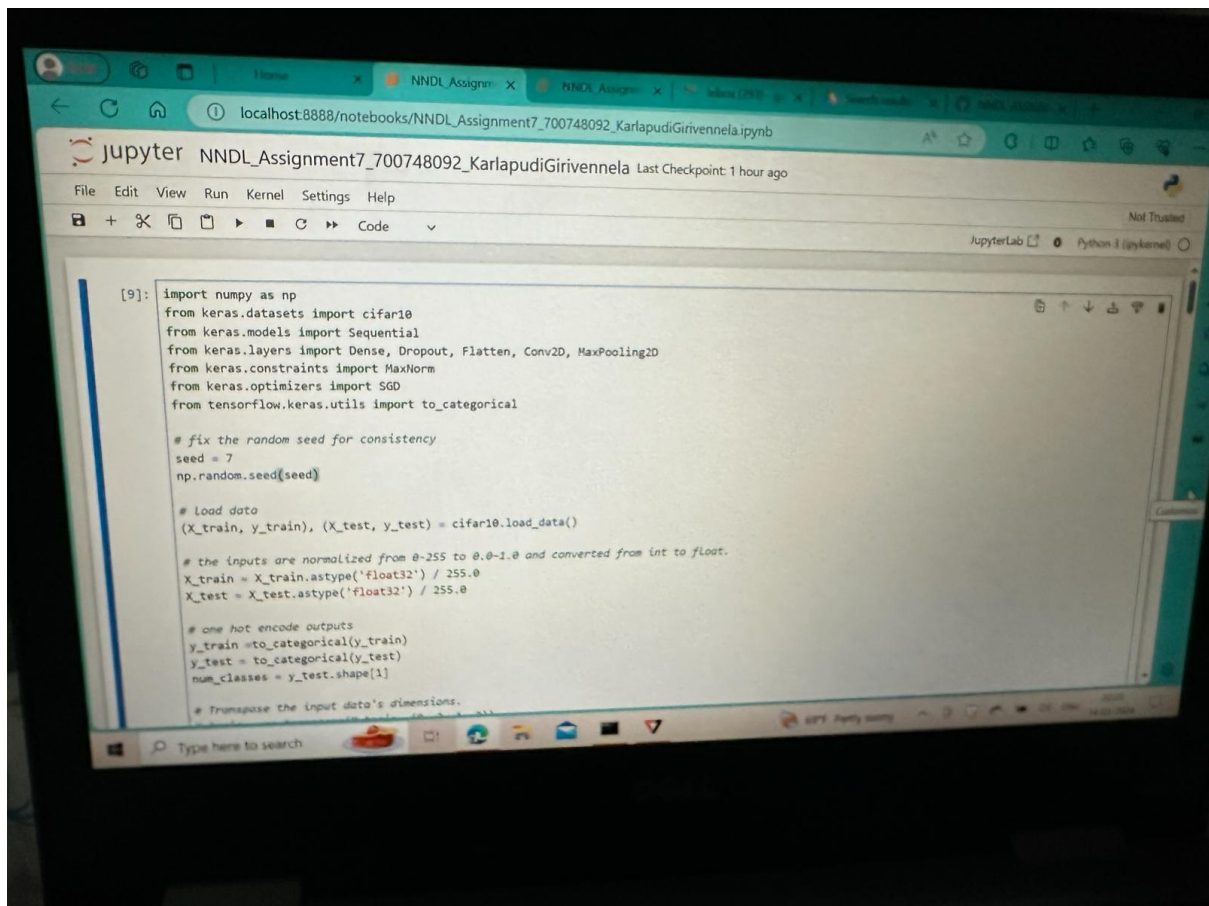
Programming elements:

1. About CNN
2. Hyperparameters of CNN
3. Image classification with CNN
4. Follow the instruction below and then report how the performance changed.(apply all at once)
 - Convolutional input layer, 32 feature maps with a size of 3×3 and a rectifier activation function. • Dropout layer at 20%. • Convolutional layer, 32 feature maps with a size of 3×3 and a rectifier activation function. • Max Pool layer with size 2×2 . • Convolutional layer, 64 feature maps with a size of 3×3 and a rectifier activation function. •

Dropout layer at 20%. • Convolutional layer, 64 feature maps with a size of 3×3 and a rectifier activation function. • Max Pool layer with size 2×2 . • Convolutional layer, 128 feature maps with a size of 3×3 and a rectifier activation function. • Dropout layer at 20%. • Convolutional layer, 128 feature maps with a size of 3×3 and a rectifier activation function. • Max Pool layer with size 2×2 . • Flatten layer. • Dropout layer at 20%. • Fully connected layer with 1024 units and a rectifier activation function. • Dropout layer at 20%. • Fully connected layer with 512 units and a rectifier activation function. • Dropout layer at 20%. • Fully connected output layer with 10 units and a Softmax activation function

GITHUB LINK:

<https://github.com/Jyoshna200/jyoshna-nueral-ass-7>



```

jupyter NNDL_Assignment7_700748092_KarlapudiGirivennela Last Checkpoint: 1 hour ago
File Edit View Run Kernel Settings Help
+ X Copy Paste Undo Redo Code

# Transpose the input data's dimensions.
X_train = np.transpose(X_train, (0, 3, 1, 2))
X_test = np.transpose(X_test, (0, 3, 1, 2))

# Create the model
model = Sequential()
model.add(Conv2D(32, (3, 3), input_shape=(3, 32, 32), padding='same', activation='relu', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(32, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Flatten())
model.add(Dense(512, activation='relu', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

# Compile model
epochs = 25
lr = 0.01
decay = lr/epochs
sgd = SGD(lr=lr, momentum=0.9, nesterov=False)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])
print(model.summary())
# Fit the model
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])

# A final assessment of the model

```

```

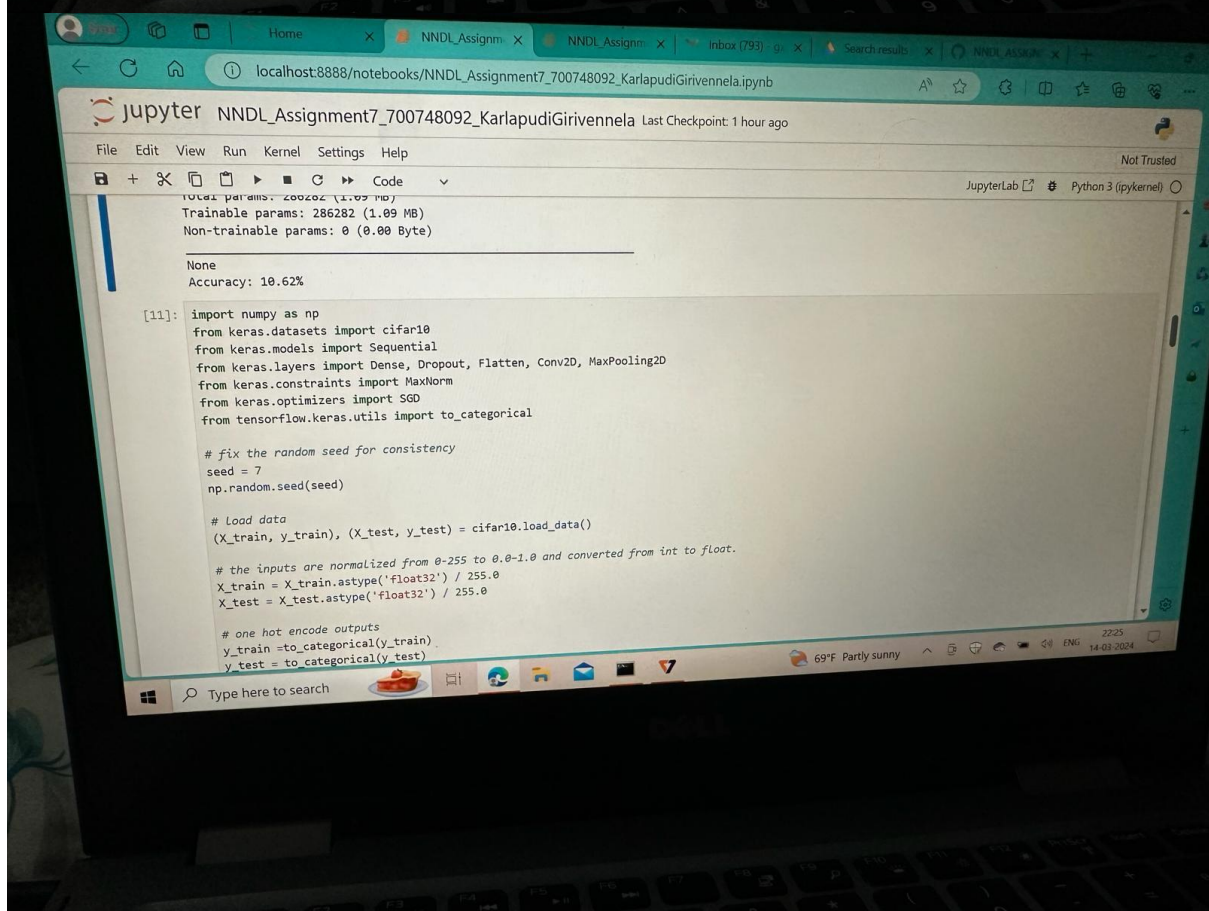
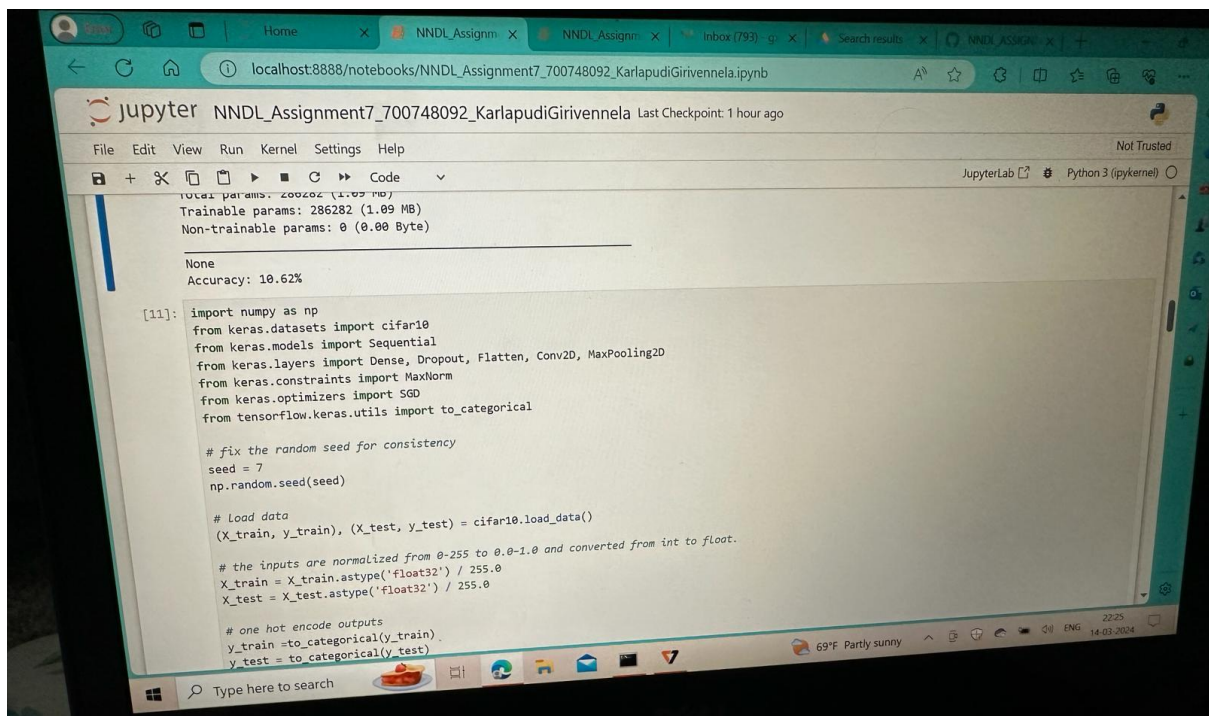
jupyter NNDL_Assignment7_700748092_KarlapudiGirivennela Last Checkpoint: 1 hour ago
File Edit View Run Kernel Settings Help
+ X Copy Paste Undo Redo Code

# A final assessment of the model.
scores = model.evaluate(X_test, y_test, verbose=1)
print("Accuracy: % .2f%%" % (scores[1]*100))

WARNING: `sgd` is deprecated as keras optimizer, please use `LearningRateScheduler` or use the legacy optimizers, e.g., of keras.optimizers.LegacyOptimizer.
Model: "sequential_3"

```

| Layer (Type) | Output Shape | Param # |
|--------------------------------|------------------|---------|
| conv2d_3 (Conv2D) | (None, 8, 8, 32) | 320 |
| dropout_3 (Dropout) | (None, 8, 8, 32) | 0 |
| conv2d_4 (Conv2D) | (None, 8, 8, 32) | 320 |
| max_pooling2d_3 (MaxPooling2D) | (None, 4, 4, 32) | 0 |
| flatten_3 (Flatten) | (None, 512) | 0 |
| dense_3 (Dense) | (None, 512) | 262016 |
| dropout_4 (Dropout) | (None, 512) | 0 |
| dense_4 (Dense) | (None, 10) | 5130 |



```

# the inputs are normalized from 0-255 to 0.0-1.0 and converted from int to float.
X_train = X_train.astype('float32') / 255.0
X_test = X_test.astype('float32') / 255.0

# one hot encode outputs
y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
num_classes = y_test.shape[1]

# Transpose the input data's dimensions.
X_train = np.transpose(X_train, (0, 3, 1, 2))
X_test = np.transpose(X_test, (0, 3, 1, 2))

# Create the model
model = Sequential()
model.add(Conv2D(32, (3, 3), input_shape=(3, 32, 32), padding='same', activation='relu', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(32, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(64, (3, 3), padding='same', activation='relu', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(MaxPooling2D(pool_size=(1, 1)))
model.add(Conv2D(128, (3, 3), padding='same', activation='relu', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(MaxPooling2D(pool_size=(1, 1)))

```

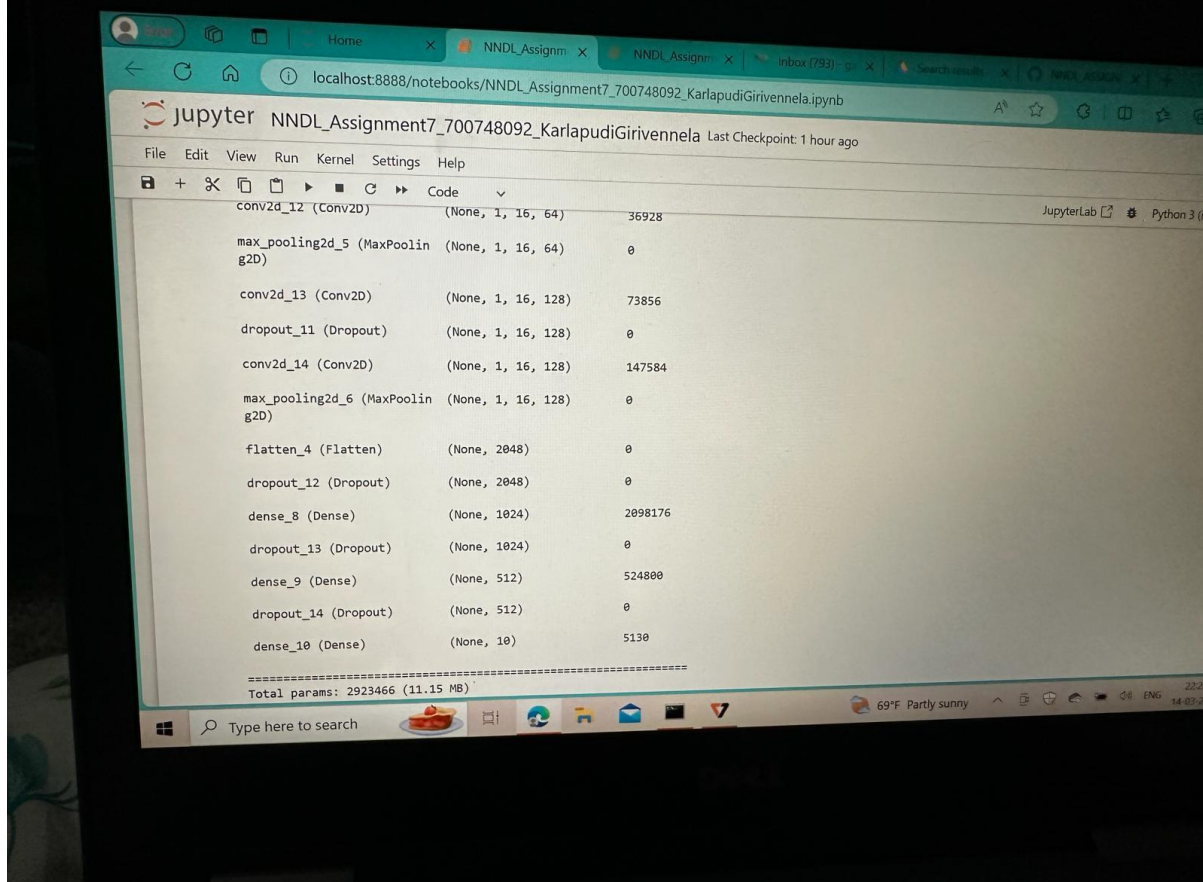
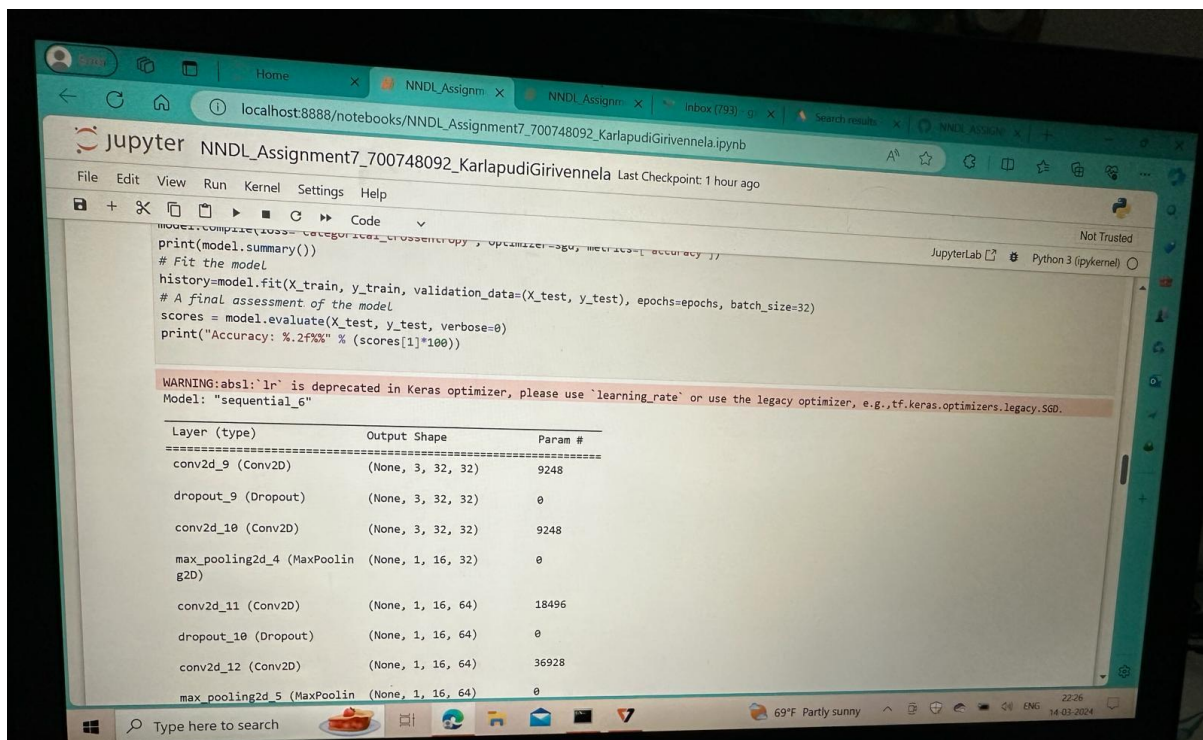
```

model.add(MaxPooling2D(pool_size=(1, 1)))
model.add(Conv2D(128, (3, 3), padding='same', activation='relu', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(MaxPooling2D(pool_size=(1, 1)))
model.add(Flatten())
model.add(Dropout(0.2))
model.add(Dense(1024, activation='relu', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))

# Compile model
epochs = 25
lr_rate = 0.01
decay = lr_rate/epochs
sgd = SGD(lr=lr_rate, momentum=0.9, nesterov=False)

model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])
print(model.summary())
# Fit the model
history=model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=32)
# A final assessment of the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1]*100))

```

```
jupyter NNDL_Assignment7_700748092_KarlapudiGirivennela.ipynb
File Edit View Run Kernel Settings Help
+ % 
Code

1563/1563 [=====] - 148s 95ms/step - loss: 1.2146 - accuracy: 0.5436 - val_loss: 1.2122 - val_accuracy: 0.5466
Epoch 14/25
1563/1563 [=====] - 147s 94ms/step - loss: 1.1936 - accuracy: 0.5485 - val_loss: 1.1993 - val_accuracy: 0.5457
Epoch 15/25
1563/1563 [=====] - 146s 94ms/step - loss: 1.1866 - accuracy: 0.5723 - val_loss: 1.1980 - val_accuracy: 0.5760
Epoch 16/25
1563/1563 [=====] - 147s 94ms/step - loss: 1.1684 - accuracy: 0.5793 - val_loss: 1.1962 - val_accuracy: 0.5760
Epoch 17/25
1563/1563 [=====] - 149s 95ms/step - loss: 1.1511 - accuracy: 0.5852 - val_loss: 1.1980 - val_accuracy: 0.5870
Epoch 18/25
1563/1563 [=====] - 146s 94ms/step - loss: 1.1454 - accuracy: 0.5891 - val_loss: 1.1824 - val_accuracy: 0.5890
Epoch 19/25
1563/1563 [=====] - 147s 94ms/step - loss: 1.1399 - accuracy: 0.5877 - val_loss: 1.1273 - val_accuracy: 0.5780
Epoch 20/25
1563/1563 [=====] - 149s 95ms/step - loss: 1.1252 - accuracy: 0.5968 - val_loss: 1.1951 - val_accuracy: 0.5870
Epoch 21/25
1563/1563 [=====] - 148s 94ms/step - loss: 1.1223 - accuracy: 0.5968 - val_loss: 1.1578 - val_accuracy: 0.5890
Epoch 22/25
1563/1563 [=====] - 154s 98ms/step - loss: 1.1093 - accuracy: 0.6068 - val_loss: 1.1799 - val_accuracy: 0.5890
Epoch 23/25
1563/1563 [=====] - 150s 96ms/step - loss: 1.1076 - accuracy: 0.6024 - val_loss: 1.1777 - val_accuracy: 0.5890
Epoch 24/25
1563/1563 [=====] - 148s 94ms/step - loss: 1.1011 - accuracy: 0.6033 - val_loss: 1.1848 - val_accuracy: 0.5890
Epoch 25/25
1563/1563 [=====] - 149s 95ms/step - loss: 1.0957 - accuracy: 0.6068 - val_loss: 1.1685 - val_accuracy: 0.5890
Accuracy: 58.84%

[12]: predictions = model.predict(X_test[:4])
print(predictions)
solution.com/nndlassignment-solution
```

```
jupyter NNDL_Assignment7_700748092_KarlapudiGirivennela.ipynb
File Edit View Run Kernel Settings Help
+ % 
Code

[12]: predictions = model.predict(X_test[:4])
print(predictions)
print(np.argmax(predictions, axis=1))
print(y_test[:4])

1/1 [=====] - 0s 182ms/step
[[6.6421755e-02 4.8251629e-02 4.9000204e-02 3.7323502e-01 8.7610923e-02
 2.3808683e-01 3.8983859e-02 2.4645276e-02 4.1188173e-02 3.2576293e-02
 4.0154178e-02 9.1406219e-02 1.3929103e-03 1.1444984e-03 9.6673996e-04
 1.9689640e-03 5.0818251e-04 1.2016086e-03 8.4895992e-01 1.2296793e-02
 5.3523266e-01 3.0359907e-02 3.6236208e-02 1.0163108e-02 1.4053379e-02
 1.1074324e-02 5.2310782e-04 6.3204574e-03 3.2090539e-01 3.5131443e-02
 7.7109390e-01 1.6146995e-02 1.9118270e-02 1.4980440e-02 4.4592754e-03
 5.2235037e-04 4.0119587e-04 1.6665384e-01 2.7205385e-03 3.9031005e-03]]
[3 8 0 0]
[[0. 0. 0. 1. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 1. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 1. 0.]
 [1. 0. 0. 0. 0. 0. 0. 0. 0. 0.]]

[13]: import matplotlib.pyplot as plt

# Plot the accuracy values for training & validation.
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
```